

Introduction to TensorFlow

Fully Connected Deep Networks with TensorFlow

郭耀仁

大綱

- 關於 Fully Connected Deep Networks
- 取得資料
- Benchmark
- 建構 TensorFlow 計算圖形
- 訓練
- 隨堂練習

關於 Fully Connected Deep Networks

Fully Connected Deep Networks 多層感知器

單一個感知器（Perceptron）的公式：

$$y = \sigma(WX + b)$$

σ 即所謂的激活函數（activation function）

激活函數用來將線性關係映射至非線性關係

常用的激活函數：


- Sigmoid
- TanH
- ReLU

連結多個感知器建構出多層感知器，透過 propagations 求出係數

- 輸入層
- 隱藏層
- 輸出層

什麼是 propagations（傳播）？

- Forward Propagation: 第一個隨機猜測
- Backward Propagation: 訓練



Source: <https://math.stackexchange.com/questions/2048722/a-name-for-layered-directed-graph-as-in-a-fully-connected-neural-network>
(<https://math.stackexchange.com/questions/2048722/a-name-for-layered-directed-graph-as-in-a-fully-connected-neural-network>).

動手做一個 Fully Connected Deep Networks

先做 Forward Propagation

```
In [1]: import numpy as np

class Neural_Network:
    def __init__(self, input_size, hidden_size, output_size):
        self._input_size = input_size
        self._hidden_size = hidden_size
        self._output_size = output_size
        self.W0 = np.random.randn(self._input_size, self._hidden_size)
        self.W1 = np.random.randn(self._hidden_size, self._output_size)
    def forward(self, X):
        self.z0 = np.dot(X, self.W0)
        self.z1 = self.sigmoid(self.z0)
        self.z2 = np.dot(self.z1, self.W1)
        output = self.sigmoid(self.z2)
        return output
    def sigmoid(self, x):
        return 1/(1 + np.exp(-x))
```

給定的 X 與 y

- 神經網路重視資料的標準化 (Standard Scaling)
 - 使用了 Activation Function 的緣故
 - Standard Scaler
 - MinMax Scaler

```
In [2]: import numpy as np

def min_max_scaler(x):
    x_min = x.min(axis=0)
    x_max = x.max(axis=0)
    return (x - x_min) / (x_max - x_min)
```

```
In [3]: X = np.linspace(1, 30, 30).reshape(-1, 1)
y = 2*X + 3
print(X.shape)
print(y.shape)
print("=====")
print(X[:5])
print("=====")
print(y[:5])
```

```
(30, 1)
```

```
(30, 1)
```

```
=====
```

```
[[1.]
```

```
 [2.]
```

```
 [3.]
```

```
 [4.]
```

```
 [5.]]
```

```
=====
```

```
[[ 5.]
```

```
 [ 7.]
```

```
 [ 9.]
```

```
[11.]
```

```
[13.]]
```

```
In [4]: X_scaled = min_max_scaler(X)
y_scaled = min_max_scaler(y)
print(X_scaled[:5])
print("=====")
print(y_scaled[:5])
```

```
[[0.
  [0.03448276]
  [0.06896552]
  [0.10344828]
  [0.13793103]]
```

```
=====
```

```
[[0.
  [0.03448276]
  [0.06896552]
  [0.10344828]
  [0.13793103]]
```

```
In [5]: NN = Neural_Network(1, 4, 1)
y_hat = NN.forward(X_scaled)
```

```
print("y_hat:")
print(y_hat.ravel())
print("=====")
print("y:")
print(y_scaled.ravel())
```

```
y_hat:
[0.4742546  0.47088419 0.46752178 0.46417296 0.46084327 0.45753813
 0.45426283 0.45102252 0.44782214 0.44466646 0.44156002 0.43850709
 0.43551173 0.4325777  0.42970851 0.42690734 0.42417713 0.42152049
 0.41893976 0.41643696 0.41401385 0.41167187 0.40941222 0.40723581
 0.4051433  0.40313508 0.40121133 0.39937198 0.39761678 0.39594525]
=====
y:
[0.          0.03448276 0.06896552 0.10344828 0.13793103 0.17241379
 0.20689655 0.24137931 0.27586207 0.31034483 0.34482759 0.37931034
 0.4137931  0.44827586 0.48275862 0.51724138 0.55172414 0.5862069
 0.62068966 0.65517241 0.68965517 0.72413793 0.75862069 0.79310345
 0.82758621 0.86206897 0.89655172 0.93103448 0.96551724 1.          ]
```


再做 Back Propagation

- 計算輸出層的 Loss
- 將 Loss 輸入 sigmoid 導函數，計算 W_1 的 Loss 比例
- W_1 的 Loss 再輸入 sigmoid 導函數，計算 W_0 的 Loss 比例
- 依據梯度遞減調整 W_1 與 W_0

```
In [6]: import numpy as np

class Neural_Network:
    def __init__(self, input_size, hidden_size, output_size):
        self._input_size = input_size
        self._hidden_size = hidden_size
        self._output_size = output_size
        self.W0 = np.random.randn(self._input_size, self._hidden_size)
        self.W1 = np.random.randn(self._hidden_size, self._output_size)
    def forward(self, X):
        self.z0 = np.dot(X, self.W0)
        self.z1 = self.sigmoid(self.z0)
        self.z2 = np.dot(self.z1, self.W1)
        output = self.sigmoid(self.z2)
        return output
    def sigmoid(self, x):
        return 1/(1 + np.exp(-x))
    def sigmoid_derivative(self, x):
        return x*(1-x)
    def backward(self, X, y, output):
        self.error = y - output
        self.output_delta = self.error * self.sigmoid_derivative(output)
        self.z1_error = np.dot(self.output_delta, self.W1.T)
        self.z1_delta = self.z1_error * self.sigmoid_derivative(self.z1)
        self.W0 += np.dot(X.T, self.z1_delta)
        self.W1 += np.dot(self.z1.T, self.output_delta)
    def train(self, X, y):
        output = self.forward(X)
        self.backward(X, y, output)
    def get_weights(self):
        return self.W0, self.W1
    def predict(self, X_test):
        return self.forward(X_test)
```

```
In [7]: NN = Neural_Network(1, 4, 1)
        for i in range(10000):
            for x_, y_ in zip(X_scaled, y_scaled):
                NN.train(x_.reshape(-1, 1), y_.reshape(-1, 1))
            loss = np.mean(((y_scaled - NN.forward(X_scaled))**2).sum())
            if i % 1000 == 0:
                print("Epoch: {}, Loss: {}".format(i, loss))
```

```
Epoch: 0, Loss: 3.2321717227043365
Epoch: 1000, Loss: 0.03812162578616321
Epoch: 2000, Loss: 0.036089745001124524
Epoch: 3000, Loss: 0.03529837545376917
Epoch: 4000, Loss: 0.03490766381521602
Epoch: 5000, Loss: 0.034702067612926185
Epoch: 6000, Loss: 0.03459881476904855
Epoch: 7000, Loss: 0.03455858545240497
Epoch: 8000, Loss: 0.03456000817600267
Epoch: 9000, Loss: 0.03459032993787807
```

```
In [8]: weight_0, weight_1 = NN.get_weights()
print(weight_0)
print(weight_1)
```

```
[[ -0.89827987  0.90333068 -0.89808679 -0.9035581  ]]
[[ -4.30576809]
 [  9.52137165]
 [ -3.89353405]
 [ -6.67238022]]
```

```
In [9]: X_test = X_scaled[:3]
y_hat_scaled = NN.predict(X_test)
y_hat = y.min() + (y.max() - y.min())*y_hat_scaled
print(y_hat)
print(y[:3])
```

```
[[ 8.73841305]
 [ 9.45883124]
 [10.30406586]]
[[5.]
 [7.]
 [9.]]
```

取得資料

簡單、作為測試目的即可

Scikit-Learn Breast Cancer 資料集

```
In [10]: from sklearn.datasets import load_breast_cancer
```

```
breast_cancer = load_breast_cancer()
print(breast_cancer.feature_names)
print(breast_cancer.DESCR)
```

```
['mean radius' 'mean texture' 'mean perimeter' 'mean area'
 'mean smoothness' 'mean compactness' 'mean concavity'
 'mean concave points' 'mean symmetry' 'mean fractal dimension'
 'radius error' 'texture error' 'perimeter error' 'area error'
 'smoothness error' 'compactness error' 'concavity error'
 'concave points error' 'symmetry error' 'fractal dimension error'
 'worst radius' 'worst texture' 'worst perimeter' 'worst area'
 'worst smoothness' 'worst compactness' 'worst concavity'
 'worst concave points' 'worst symmetry' 'worst fractal dimension']
.. _breast_cancer_dataset:
```

Breast cancer wisconsin (diagnostic) dataset

****Data Set Characteristics:****

:Number of Instances: 569

:Number of Attributes: 30 numeric, predictive attributes and the class

:Attribute Information:

- radius (mean of distances from center to points on the perimeter)
- texture (standard deviation of gray-scale values)
- perimeter
- area
- smoothness (local variation in radius lengths)
- compactness ($\text{perimeter}^2 / \text{area} - 1.0$)
- concavity (severity of concave portions of the contour)
- concave points (number of concave portions of the contour)
- symmetry

- fractal dimension ("coastline approximation" - 1)

The mean, standard error, and "worst" or largest (mean of the three largest values) of these features were computed for each image, resulting in 30 features. For instance, field 3 is Mean Radius, field 13 is Radius SE, field 23 is Worst Radius.

- class:
 - WDBC-Malignant
 - WDBC-Benign

:Summary Statistics:

| ===== | ===== | ===== |
|-------------------------------------|-------|--------|
| | Min | Max |
| ===== | ===== | ===== |
| radius (mean): | 6.981 | 28.11 |
| texture (mean): | 9.71 | 39.28 |
| perimeter (mean): | 43.79 | 188.5 |
| area (mean): | 143.5 | 2501.0 |
| smoothness (mean): | 0.053 | 0.163 |
| compactness (mean): | 0.019 | 0.345 |
| concavity (mean): | 0.0 | 0.427 |
| concave points (mean): | 0.0 | 0.201 |
| symmetry (mean): | 0.106 | 0.304 |
| fractal dimension (mean): | 0.05 | 0.097 |
| radius (standard error): | 0.112 | 2.873 |
| texture (standard error): | 0.36 | 4.885 |
| perimeter (standard error): | 0.757 | 21.98 |
| area (standard error): | 6.802 | 542.2 |
| smoothness (standard error): | 0.002 | 0.031 |
| compactness (standard error): | 0.002 | 0.135 |
| concavity (standard error): | 0.0 | 0.396 |
| concave points (standard error): | 0.0 | 0.053 |
| symmetry (standard error): | 0.008 | 0.079 |
| fractal dimension (standard error): | 0.001 | 0.03 |
| radius (worst): | 7.93 | 36.04 |
| texture (worst): | 12.02 | 49.54 |

| | | |
|----------------------------|-------|--------|
| perimeter (worst): | 50.41 | 251.2 |
| area (worst): | 185.2 | 4254.0 |
| smoothness (worst): | 0.071 | 0.223 |
| compactness (worst): | 0.027 | 1.058 |
| concavity (worst): | 0.0 | 1.252 |
| concave points (worst): | 0.0 | 0.291 |
| symmetry (worst): | 0.156 | 0.664 |
| fractal dimension (worst): | 0.055 | 0.208 |
| ===== | ===== | ===== |

:Missing Attribute Values: None

:Class Distribution: 212 - Malignant, 357 - Benign

:Creator: Dr. William H. Wolberg, W. Nick Street, Olvi L. Mangasarian

:Donor: Nick Street

:Date: November, 1995

This is a copy of UCI ML Breast Cancer Wisconsin (Diagnostic) datasets.
<https://goo.gl/U2Uwz2>

Features are computed from a digitized image of a fine needle aspirate (FNA) of a breast mass. They describe characteristics of the cell nuclei present in the image.

Separating plane described above was obtained using Multisurface Method-Tree (MSM-T) [K. P. Bennett, "Decision Tree Construction Via Linear Programming." Proceedings of the 4th Midwest Artificial Intelligence and Cognitive Science Society, pp. 97-101, 1992], a classification method which uses linear programming to construct a decision tree. Relevant features were selected using an exhaustive search in the space of 1-4 features and 1-3 separating planes.

The actual linear program used to obtain the separating plane in the 3-dimensional space is that described in:

[K. P. Bennett and O. L. Mangasarian: "Robust Linear Programming Discrimination of Two Linearly Inseparable Sets", Optimization Methods and Software 1, 1992, 23-34].

This database is also available through the UW CS ftp server:

```
ftp ftp.cs.wisc.edu
cd math-prog/cpo-dataset/machine-learn/WDBC/
```

.. topic:: References

- W.N. Street, W.H. Wolberg and O.L. Mangasarian. Nuclear feature extraction
for breast tumor diagnosis. IS&T/SPIE 1993 International Symposium on
Electronic Imaging: Science and Technology, volume 1905, pages 861-870,
San Jose, CA, 1993.
- O.L. Mangasarian, W.N. Street and W.H. Wolberg. Breast cancer diagnosis and
prognosis via linear programming. Operations Research, 43(4), pages 570-5
77,
July-August 1995.
- W.H. Wolberg, W.N. Street, and O.L. Mangasarian. Machine learning techniques
to diagnose breast cancer from fine-needle aspirates. Cancer Letters 77
(1994)
163-171.

```
In [11]: X_arr = breast_cancer.data  
y_arr = breast_cancer.target  
print(X_arr.shape)  
print(y_arr.shape)
```

```
(569, 30)
```

```
(569,)
```

```
In [12]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_arr, y_arr, test_size=0.3, r
andom_state=123)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(398, 30)
(171, 30)
(398,)
(171,)
```

Benchmark

```
In [13]: from sklearn.linear_model import LogisticRegression
         from sklearn.metrics import accuracy_score

         clf = LogisticRegression(solver="liblinear")
         clf.fit(X_train, y_train)
         y_pred = clf.predict(X_test)
         acc = accuracy_score(y_test, y_pred)
```

```
In [14]: print(clf.coef_)  
print(clf.intercept_)  
print(acc)
```

```
[[ 1.93788057e+00  1.94992830e-01 -2.93534745e-03 -9.97160018e-03  
 -1.31053304e-01 -3.94404828e-01 -5.87789964e-01 -3.08240545e-01  
 -2.34487470e-01 -2.08028628e-02 -1.37757133e-03  1.25279091e+00  
  8.06843035e-03 -8.58257925e-02 -1.44335700e-02 -4.29282110e-04  
 -4.53672393e-02 -3.74242144e-02 -3.87957326e-02  6.72139227e-03  
  1.00102118e+00 -3.88323187e-01 -1.28745720e-01 -1.78899170e-02  
 -2.50197148e-01 -1.09822573e+00 -1.45824679e+00 -6.35170781e-01  
 -7.02595973e-01 -9.59181366e-02]]  
[0.41185323]  
0.9824561403508771
```

Benchmark 完整程式碼

```
In [15]: from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

breast_cancer = load_breast_cancer()
X_arr = breast_cancer.data
y_arr = breast_cancer.target
X_train, X_test, y_train, y_test = train_test_split(X_arr, y_arr, test_size=0.3, r
andom_state=123)
clf = LogisticRegression(solver="liblinear")
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
acc = accuracy_score(y_test, y_pred)
```


建構 TensorFlow 計算圖形

準備 Placeholders 作為 Input layers

```
In [16]: import tensorflow as tf

n_features = X_train.shape[1]
with tf.name_scope("input-layer"):
    X = tf.placeholder(tf.float32)
    y = tf.placeholder(tf.float32)
```

Placeholders 不指定外觀的好處

- 多層感知器的應用常會使用 **Mini-batching** 的技巧
- 訓練結束之後可以餵入 `X_test` 獲得 `y_pred_arr`

準備 Variables 作為 Hidden layers

In [17]: `import numpy as np`

```
n_classes = np.unique(y_train).size
n_neurons = 2**4
with tf.name_scope("hidden-layer"):
    W = tf.Variable(tf.random_normal((n_features, n_neurons))) # (30, 16)
    b = tf.Variable(tf.random_normal((n_neurons,))) # (16,)
    X_hidden = tf.nn.relu(tf.matmul(X, W) + b) # (398, 30) x (30,
16) + (16,)
```

WARNING:tensorflow:From /Users/kuoyaojen/anaconda3/envs/tensorflow/lib/python3.6/site-packages/tensorflow/python/framework/op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

如何決定 hidden-layer 的層數與神經元數量？

- 絕大多數情境：使用一層 hidden-layer 就足夠
- 神經元數量： 2^n

準備 Variables 作為 Output layers

```
In [18]: with tf.name_scope("output-layer"):
          W = tf.Variable(tf.random_normal((n_neurons, 1))) # (16, 1)
          b = tf.Variable(tf.random_normal((1,)))           # (1,)
          y_logit = tf.squeeze(tf.add(tf.matmul(X_hidden, W), b)) # (398, 16) x (16, 1)
          + (1,)
          y_one_prob = tf.sigmoid(y_logit)
          y_pred = tf.round(y_one_prob)
```

寫下成本函數的公式

```
In [19]: with tf.name_scope("loss"):  
          entropy = tf.nn.sigmoid_cross_entropy_with_logits(logits=y_logit, labels=y)  
          loss = tf.reduce_sum(entropy)
```


宣告 Optimizer 與學習速率

```
In [20]: learning_rate = 0.0001
         with tf.name_scope("optimizer"):
             optimizer = tf.train.AdamOptimizer(learning_rate).minimize(loss)
```

WARNING:tensorflow:From /Users/kuoyaojen/anaconda3/envs/tensorflow/lib/python3.6/site-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:
Use tf.cast instead.

建構 TensorFlow 計算圖形完整程式碼

```
In [21]: import tensorflow as tf
import numpy as np

tf.reset_default_graph()
n_features = X_train.shape[1]
n_classes = np.unique(y_train).size
n_neurons = 2**4
learning_rate = 0.0001

with tf.name_scope("input-layer"):
    X = tf.placeholder(tf.float32)
    y = tf.placeholder(tf.float32)
with tf.name_scope("hidden-layer"):
    W = tf.Variable(tf.random_normal((n_features, n_neurons))) # (30, 16)
    b = tf.Variable(tf.random_normal((n_neurons,))) # (16,)
    X_hidden = tf.nn.relu(tf.matmul(X, W) + b) # (398, 30) x (30,
    16) + (16,)
with tf.name_scope("output-layer"):
    W = tf.Variable(tf.random_normal((n_neurons, 1))) # (16, 1)
    b = tf.Variable(tf.random_normal((1,))) # (1,)
    y_logit = tf.squeeze(tf.add(tf.matmul(X_hidden, W), b)) # (398, 16) x (16, 1)
    + (1,)
    y_one_prob = tf.sigmoid(y_logit)
    y_pred = tf.round(y_one_prob)
with tf.name_scope("loss"):
    entropy = tf.nn.sigmoid_cross_entropy_with_logits(logits=y_logit, labels=y)
    loss = tf.reduce_sum(entropy)
with tf.name_scope("optimizer"):
    optimizer = tf.train.AdamOptimizer(learning_rate).minimize(loss)
```

訓練

```
In [22]: n_steps = 50000
file_writer_path = "./graphs/fully-connected-deep-networks"
loss_history = []

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer()) # 初始化所有的變數張量!
    train_writer = tf.summary.FileWriter(file_writer_path, tf.get_default_graph())
    for i in range(n_steps):
        feed_dict = {
            X: X_train,
            y: y_train
        }
        _, loss_ = sess.run([optimizer, loss], feed_dict=feed_dict)
        loss_history.append(loss_)
        if i % 1000 == 0:
            print("step {}, loss: {}".format(i, loss_))
    w_final, b_final = sess.run([W, b])
    y_pred_arr = sess.run(y_pred, feed_dict={X: X_test})
```

```
step 0, loss: 539639.75
step 1000, loss: 274796.96875
step 2000, loss: 37507.2578125
step 3000, loss: 5371.13427734375
step 4000, loss: 5262.8720703125
step 5000, loss: 5168.35546875
step 6000, loss: 5027.861328125
step 7000, loss: 4896.12841796875
step 8000, loss: 3382.2109375
step 9000, loss: 2716.5947265625
step 10000, loss: 2126.1630859375
step 11000, loss: 1527.1834716796875
step 12000, loss: 977.72314453125
step 13000, loss: 590.98681640625
step 14000, loss: 299.52642822265625
step 15000, loss: 104.7797622680664
step 16000, loss: 68.5936279296875
step 17000, loss: 59.4394416809082
```

```
In [23]: import matplotlib.pyplot as plt

plt.plot(range(n_steps), loss_history)
plt.title("Loss Summary")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.show()
```

<Figure size 640x480 with 1 Axes>

```
In [24]: acc = accuracy_score(y_test, y_pred_arr)
print(w_final)
print(b_final)
print(acc)
```

```
[[ 0.26682755]
 [-0.53566235]
 [-0.10864915]
 [ 1.4385455 ]
 [-0.6868179 ]
 [-1.319403  ]
 [ 0.5819109 ]
 [ 0.06051218]
 [ 0.2184967 ]
 [ 0.6089343 ]
 [-0.143741  ]
 [-1.5354853 ]
 [-0.6753923 ]
 [-1.8530135 ]
 [ 1.4804659 ]
 [ 0.6578026 ]]
[2.2221074]
0.9766081871345029
```

隨堂練習

使用 titanic 資料集並利用 TensorFlow 建立一個 Fully Connected Deep Networks Classifier

```
In [1]: import pandas as pd

train_url = "https://s3-ap-northeast-1.amazonaws.com/kaggle-getting-started/titanic/train.csv"
train = pd.read_csv(train_url)
```


使用一個 Linear Regressor 填補 Age 變數的遺漏值

```
In [2]: train.describe()
```

Out[2]:

| | PassengerId | Survived | Pclass | Age | SibSp | Parch | Fare |
|-------|-------------|------------|------------|------------|------------|------------|------------|
| count | 891.000000 | 891.000000 | 891.000000 | 714.000000 | 891.000000 | 891.000000 | 891.000000 |
| mean | 446.000000 | 0.383838 | 2.308642 | 29.699118 | 0.523008 | 0.381594 | 32.204208 |
| std | 257.353842 | 0.486592 | 0.836071 | 14.526497 | 1.102743 | 0.806057 | 49.693429 |
| min | 1.000000 | 0.000000 | 1.000000 | 0.420000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 223.500000 | 0.000000 | 2.000000 | 20.125000 | 0.000000 | 0.000000 | 7.910400 |
| 50% | 446.000000 | 0.000000 | 3.000000 | 28.000000 | 0.000000 | 0.000000 | 14.454200 |
| 75% | 668.500000 | 1.000000 | 3.000000 | 38.000000 | 1.000000 | 0.000000 | 31.000000 |
| max | 891.000000 | 1.000000 | 3.000000 | 80.000000 | 8.000000 | 6.000000 | 512.329200 |

```
In [17]: from sklearn.linear_model import LinearRegression
import numpy as np

is_nan = np.isnan(train["Age"])
not_nan = ~(np.isnan(train["Age"]))
X_train = train[not_nan].loc[:, ["Pclass", "SibSp", "Parch", "Fare"]].values
y_train = train[not_nan]["Age"].values
regressor = LinearRegression()
regressor.fit(X_train, y_train)
X_test = train[is_nan].loc[:, ["Pclass", "SibSp", "Parch", "Fare"]].values
y_hat = regressor.predict(X_test)
```

```
In [ ]:
```